

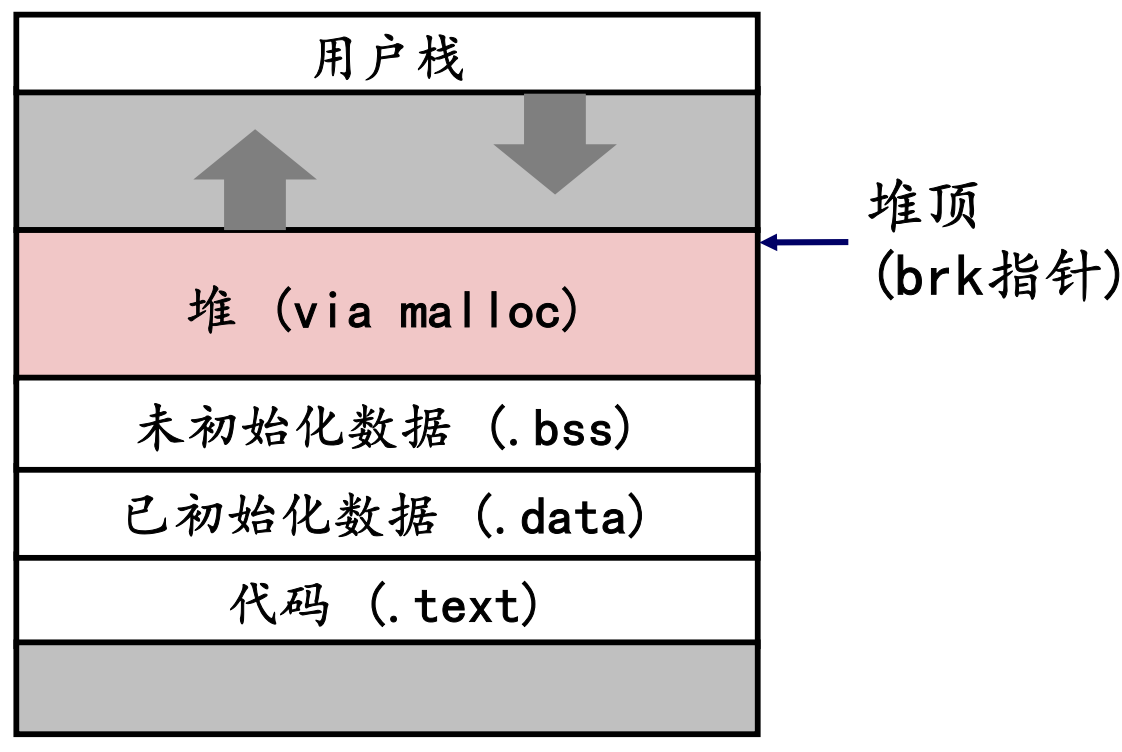

动态内存分配：基本概念 (Dynamic Memory Allocation)

主要内容

- 基本概念
- 隐式空闲列表

动态内存分配

- 程序员使用**动态内存分配器 (Dynamic memory allocator)** (例如 `malloc`) 在运行时获取 VM。
 - 对于仅在运行时才知道大小的数据结构。
- 动态内存分配器管理称为**堆 (Heap)**的进程虚拟内存区域。



动态内存分配

- 分配器将堆维护为可变大小的块^块的集合，这些块要么被分配^{分配}，要么被释放^{释放}
- 分配器的类型
 - 显式分配器：应用程序分配和释放空间
 - E. g., C语言中 malloc 和 free
 - 隐式分配器：应用程序分配空间，但不释放空间
 - 例如 Java、ML 和 Lisp 中的垃圾回收
- 今天将讨论简单的显式内存分配

malloc程序包

```
#include <stdlib.h>
```

```
void *malloc(size_t size)
```

- 成功时：
 - 返回一个指针，指向至少包含 `size` 字节的内存块，并且内存块起始地址按照 8 字节（x86）或 16 字节（x86-64）的边界对齐。
 - 如果 `size == 0`，返回 `NULL`。
- 失败时：返回 `NULL` (0) 并设置 `errno`。

```
void free(void *p)
```

- 将指针 `p` 所指向的内存块归还到可用内存池。
- 参数 `p` 必须来自于之前的 `malloc` 或 `realloc` 调用。

其他函数

- `calloc`: `malloc` 的一个版本，分配的内存会初始化为 0。
- `realloc`: 改变之前已经分配的内存块大小。
- `sbrk`: 分配器在内部调用它来扩展或收缩堆。

Malloc例子

```
#include <stdio.h>
#include <stdlib.h>

void foo(int n) {
    int i, *p;


    /* Allocate a block of n ints */
    p = (int *) malloc(n * sizeof(int));
    if (p == NULL) {
        perror("malloc");
        exit(0);
    }

    /* Initialize allocated block */
    for (i=0; i<n; i++)
        p[i] = i;

    /* Return allocated block to the heap */
    free(p);
}
```

本课堂中的假设

- 内存是按字寻址的。
- 字的大小是整型数(int)。




分配例子


```
p1 = malloc(4)
```



```
p2 = malloc(5)
```




```
p3 = malloc(6)
```



```
free(p2)
```



```
p4 = malloc(2)
```



约束条件

■ 应用程序

- 可以发出任意顺序的 `malloc` 和 `free` 请求
- `free` 请求必须针对已 `malloc` 的块

■ 分配器

- 无法控制已分配块的数量或大小
- 必须立即响应 `malloc` 请求
 - 即，无法重新排序或缓冲请求
- 必须从空闲内存中分配块
 - 即，只能将已分配的块放置在空闲内存中
- 必须对齐块以满足所有对齐要求
- 在 Linux 系统上，8 字节 (x86) 或 16 字节 (x86-64) 对齐
- 只能操作和修改空闲内存
- 分配的块一旦 `malloc` 分配完毕，就无法移动
 - 即，不允许进行压缩

性能目标：吞吐量

- 给定一些 `malloc` 和 `free` 请求序列

- $R_0, R_1, \dots, R_k, \dots, R_{n-1}$

- 目标：最大化吞吐量和峰值内存利用率

- 这些目标通常相互冲突

- 吞吐量：

- 单位时间内完成的请求数量

- 示例：

- 10 秒内 5,000 次 `malloc` 调用和 5,000 次 `free` 调用

- 吞吐量为每秒 1,000 次操作

性能目标：峰值内存利用率

- 给定一些 `malloc` 和 `free` 请求序列：
 - $R_0, R_1, \dots, R_k, \dots, R_{n-1}$
- **定义：**总有效负载 P_k
 - `malloc(p)` 会分配一个大小为 `p` 字节的块（称为 `payload`）
 - 当请求 R_k 完成后，总有效负载 P_k 等于当前所有已分配块的 `payload` 之和
- **定义：**当前堆大小 H_k
 - 假设 H_k 是单调不减的
 - 也就是说，堆大小只会在分配器调用 `sbrk` 时增长
- **定义：**第 $k+1$ 个请求后的峰值内存利用率 U_k


$$U_k = \frac{\max_{i \leq k} P_i}{H_k}$$

碎片化

- 碎片导致内存利用率低
 - 内部碎片
 - 外部碎片

内部碎片

- 对于一个给定的块，如果 有效负载 (payload) 小于块的大小，就会产生内部碎片 (internal fragmentation)。



■ 成因：

- 维护堆数据结构的开销
- 为对齐 (alignment) 而填充的字节
- 明确的策略决策
(例如：为了满足一个小请求，返回了一个较大的块)


■ 特点：

- 内部碎片 只依赖于过去的请求模式
- 因此，它相对容易测量


外部碎片

- 当堆中的总空闲内存足够大，但没有一个连续的空闲块足够大时，就会产生外部碎片。


`p1 = malloc(4)`




`p2 = malloc(5)`



`p3 = malloc(6)`



`free(p2)`



`p4 = malloc(6)`

这时会发生什么？

- 外部碎片 依赖于未来请求的模式
 - 因此，很难准确估计


Implementation Issues

- **How do we know how much memory to free given just a pointer?**
- **How do we keep track of the free blocks?**
- **What do we do with the extra space when allocating a structure that is smaller than the free block it is placed in?**
- **How do we pick a block to use for allocation -- many might fit?**
- **How do we reinsert freed block?**

Knowing How Much to Free


■ Standard method

- Keep the length of a block in the word preceding the block.
 - This word is often called the *header field* or *header*
- Requires an extra word for every allocated block




Keeping Track of Free Blocks

- Method 1: *Implicit list* using length—links all blocks



- Method 2: *Explicit list* among the free blocks using pointers



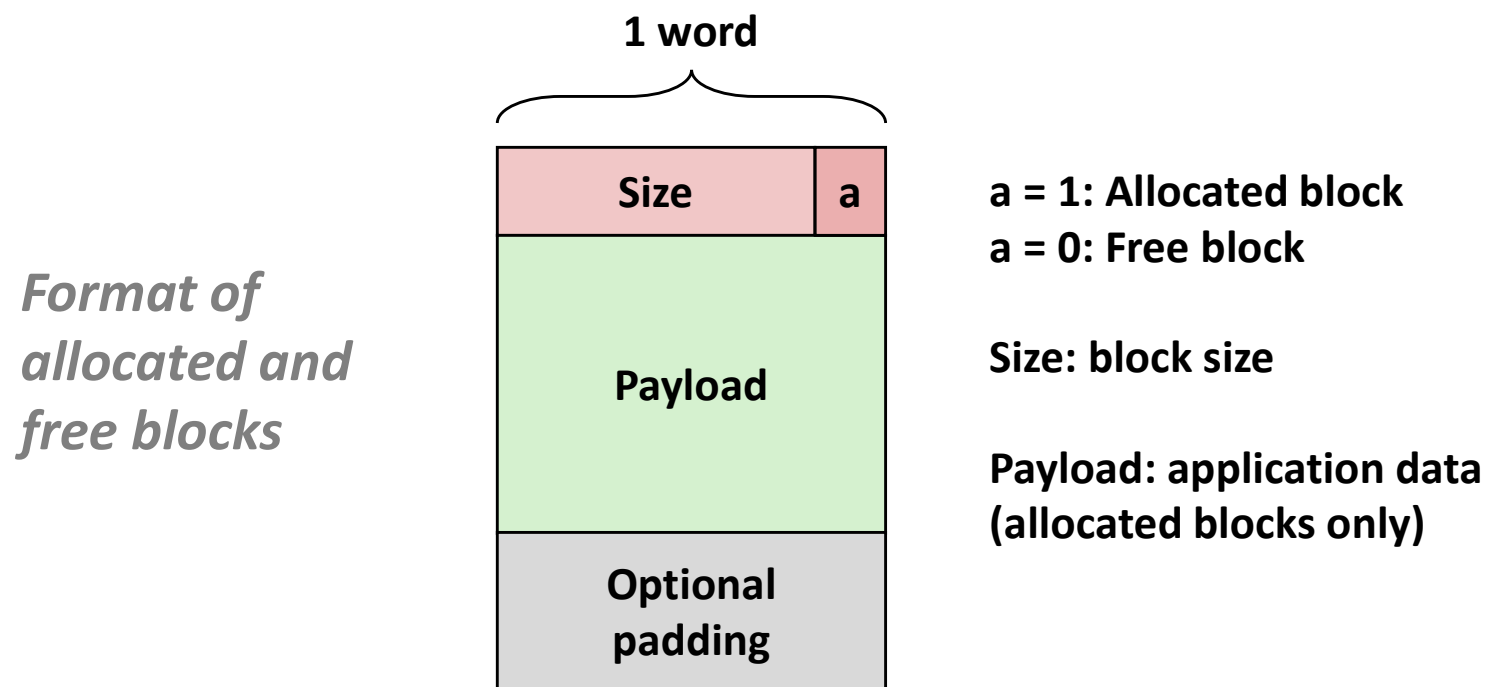
- Method 3: *Segregated free list*
 - Different free lists for different size classes
- Method 4: *Blocks sorted by size*
 - Can use a balanced tree (e.g. Red-Black tree) with pointers within each free block, and the length used as a key

Today


- Basic concepts
- **Implicit free lists**

Method 1: Implicit List

- For each block we need both size and allocation status
 - Could store this information in two words: wasteful!
- Standard trick
 - If blocks are aligned, some low-order address bits are always 0
 - Instead of storing an always-0 bit, use it as a allocated/free flag
 - When reading size word, must mask out this bit



Detailed Implicit Free List Example



Double-word
aligned

Allocated blocks: shaded

Free blocks: unshaded

Headers: labeled with size in bytes/allocated bit

Implicit List: Finding a Free Block

■ *First fit:*

- Search list from beginning, choose *first* free block that fits:

```
p = start;
while ((p < end) &&          \\ not passed end
      ((*p & 1) ||          \\ already allocated
      (*p <= len)))        \\ too small
  p = p + (*p & -2);        \\ goto next block (word addressed)
```

- Can take linear time in total number of blocks (allocated and free)
- In practice it can cause “splinters” at beginning of list

■ *Next fit:*

- Like first fit, but search list starting where previous search finished
- Should often be faster than first fit: avoids re-scanning unhelpful blocks
- Some research suggests that fragmentation is worse


■ *Best fit:*

- Search the list, choose the *best* free block: fits, with fewest bytes left over
- Keeps fragments small—usually improves memory utilization
- Will typically run slower than first fit


Implicit List: Allocating in Free Block

■ Allocating in a free block: *splitting*

- Since allocated space might be smaller than free space, we might want to split the block



`addblock(p, 4)`



```
void addblock(ptr p, int len) {
    int newsize = ((len + 1) >> 1) << 1; // round up to even
    int oldsize = *p & -2;                // mask out low bit
    *p = newsize | 1;                      // set new length
    if (newsize < oldsize)
        *(p+newsize) = oldsize - newsize; // set length in remaining
                                           // part of block
}
```


Implicit List: Freeing a Block

■ Simplest implementation:

- Need only clear the “allocated” flag

```
void free_block(ptr p) { *p = *p & -2 }
```

- But can lead to “false fragmentation”




malloc(5) ***Oops!***

There is enough free space, but the allocator won't be able to find it

Implicit List: Coalescing

- Join (*coalesce*) with next/previous blocks, if they are free
 - Coalescing with next block




```
void free_block(ptr p) {
    *p = *p & -2;           // clear allocated flag
    next = p + *p;          // find next block
    if ((*next & 1) == 0)
        *p = *p + *next;    // add to this block if
    // not allocated
}
```

- But how do we coalesce with *previous* block?


Implicit List: Bidirectional Coalescing

■ *Boundary tags* [Knuth73]

- Replicate size/allocated word at “bottom” (end) of free blocks
- Allows us to traverse the “list” backwards, but requires extra space
- Important and general technique!



*Format of
allocated and
free blocks*

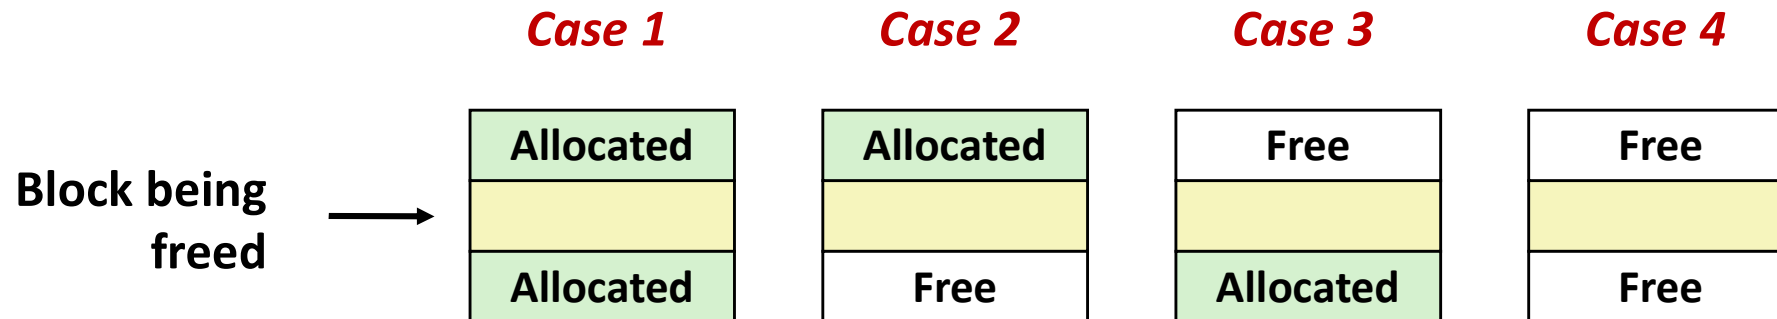


a = 1: Allocated block
a = 0: Free block


Size: Total block size

**Payload: Application data
(allocated blocks only)**

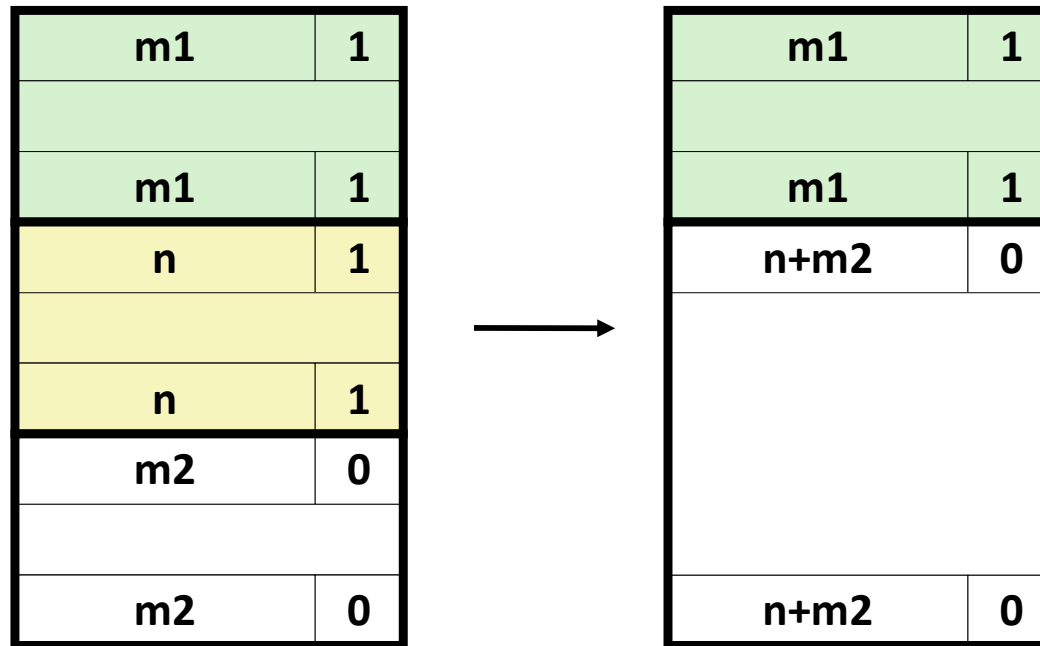
Constant Time Coalescing



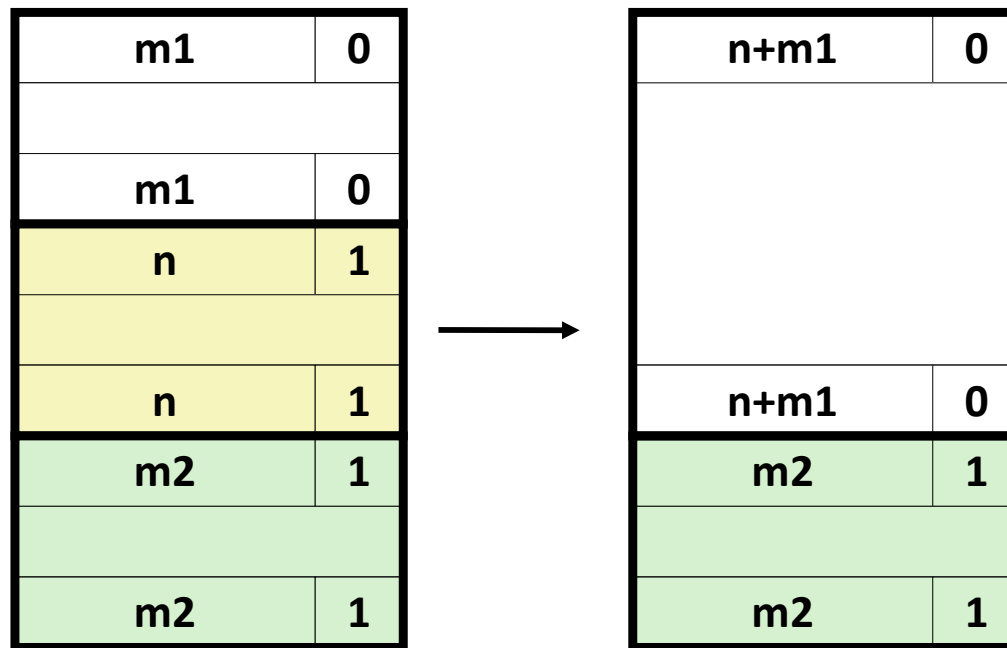
Constant Time Coalescing (Case 1)




Constant Time Coalescing (Case 2)



Constant Time Coalescing (Case 3)



Constant Time Coalescing (Case 4)



Disadvantages of Boundary Tags

- **Internal fragmentation**
- **Can it be optimized?**
 - Which blocks need the footer tag?
 - What does that mean?

Summary of Key Allocator Policies

■ Placement policy:

- First-fit, next-fit, best-fit, etc.
- Trades off lower throughput for less fragmentation
- *Interesting observation*: segregated free lists (next lecture)
approximate a best fit placement policy without having to search entire free list

■ Splitting policy:

- When do we go ahead and split free blocks?
- How much internal fragmentation are we willing to tolerate?

■ Coalescing policy:

- *Immediate coalescing*: coalesce each time **free** is called
- *Deferred coalescing*: try to improve performance of **free** by deferring coalescing until needed. Examples:
 - Coalesce as you scan the free list for **malloc**
 - Coalesce when the amount of external fragmentation reaches some threshold

Implicit Lists: Summary

- **Implementation: very simple**
- **Allocate cost:**
 - linear time worst case
- **Free cost:**
 - constant time worst case
 - even with coalescing
- **Memory usage:**
 - will depend on placement policy
 - First-fit, next-fit or best-fit
- **Not used in practice for `malloc/free` because of linear-time allocation**
 - used in many special purpose applications
- **However, the concepts of splitting and boundary tag coalescing are general to *all* allocators**